

# Assignment 1

February 23, 2024

## 1 CS6700: Reinforcement Learning

### 1.1 Programming Assignment 1

Submitted by: - Archish S (ME20B032) - Vinayak Gupta (EE20B152)

## 2 Imports

### 2.1 Libraries

```
[1]: import numpy as np
from math import floor
import matplotlib.pyplot as plt
import seaborn as sns

import tqdm
```

```
[2]: num_rows = 10
num_columns = 10
start_state = {
    "s1": np.array([[0, 4]]),
    "s2": np.array([[3, 6]]),
}
obstructions = np.array([[0, 7], [1, 1], [1, 2], [1, 3], [1, 7], [2, 1], [2, 3],
                        [2, 7], [3, 1], [3, 3], [3, 5], [4, 3], [4, 5], [4, 7],
                        [5, 3], [5, 7], [5, 9], [6, 3], [6, 9], [7, 1], [7, 6],
                        [7, 7], [7, 8], [7, 9], [8, 1], [8, 5], [8, 6], [9, 1]])

bad_states = np.array([[1, 9], [4, 2], [4, 4], [7, 5], [9, 9]])
restart_states = np.array([[3, 7], [8, 2]])
goal_states = np.array([[0, 9], [2, 2], [8, 7]])

step_reward = -1
goal_reward = 10
bad_state_reward = -6
restart_state_reward = -100

p_good_transition = 1
```

```
bias = 0.5
```

## 2.2 Environment Definition

```
[3]: class BaseEnv:
```

```
    def step(self, state, action):
        raise NotImplementedError

    def reset(self):
        raise NotImplementedError
```

```
[4]: class GridWorld(BaseEnv):
```

```
    """
    Creates a gridworld object to pass to an RL algorithm.
    Parameters
    -----
    num_rows : int
        The number of rows in the gridworld.
    num_cols : int
        The number of cols in the gridworld.
    start_state : numpy array of shape (1, 2), np.array([[row, col]])
        The start state of the gridworld (can only be one start state)
    goal_states : numpy array of shape (n, 2)
        The goal states for the gridworld where n is the number of goal
        states.
    """
    def __init__(self, num_rows, num_cols, start_state, goal_states,
↪wind=False):
        self.num_rows = num_rows
        self.num_cols = num_cols
        self.start_state = start_state
        self.goal_states = goal_states
        self.obs_states = None
        self.bad_states = None
        self.num_bad_states = 0
        self.p_good_trans = None
        self.bias = None
        self.r_step = None
        self.r_goal = None
        self.r_dead = None
        self.gamma = 1 # default is no discounting
        self.wind = wind

    def add_obstructions(self, obstructed_states=None, bad_states=None,
↪restart_states=None):
```

```

self.obs_states = obstructed_states
self.bad_states = bad_states
if bad_states is not None:
    self.num_bad_states = bad_states.shape[0]
else:
    self.num_bad_states = 0
self.restart_states = restart_states
if restart_states is not None:
    self.num_restart_states = restart_states.shape[0]
else:
    self.num_restart_states = 0

def add_transition_probability(self, p_good_transition, bias):

    self.p_good_trans = p_good_transition
    self.bias = bias

def add_rewards(self, step_reward, goal_reward, bad_state_reward=None,
↪restart_state_reward=None):

    self.r_step = step_reward
    self.r_goal = goal_reward
    self.r_bad = bad_state_reward
    self.r_restart = restart_state_reward

def create_gridworld(self):

    self.num_actions = 4
    self.num_states = self.num_cols * self.num_rows# +1
    self.start_state_seq = GridWorld.row_col_to_seq(self.start_state, self.
↪num_cols)
    self.goal_states_seq = GridWorld.row_col_to_seq(self.goal_states, self.
↪num_cols)

    # rewards structure
    self.R = self.r_step * np.ones((self.num_states, 1))
    #self.R[self.num_states-1] = 0
    self.R[self.goal_states_seq] = self.r_goal

    for i in range(self.num_bad_states):
        if self.r_bad is None:
            raise Exception("Bad state specified but no reward is given")
        bad_state = GridWorld.row_col_to_seq(self.bad_states[i,:].
↪reshape(1,-1), self.num_cols)
        #print("bad states", bad_state)
        self.R[bad_state, :] = self.r_bad

```

```

    for i in range(self.num_restart_states):
        if self.r_restart is None:
            raise Exception("Restart state specified but no reward is_
↳given")
        restart_state = GridWorld.row_col_to_seq(self.restart_states[i,:].
↳reshape(1,-1), self.num_cols)
        #print("restart_state", restart_state)
        self.R[restart_state, :] = self.r_restart

    # probability model
    if self.p_good_trans == None:
        raise Exception("Must assign probability and bias terms via the_
↳add_transition_probability method.")

    self.P = np.zeros((self.num_states, self.num_states, self.num_actions))
    for action in range(self.num_actions):
        for state in range(self.num_states):

            # check if the state is the goal state or an obstructed state -_
↳transition to end
            row_col = GridWorld.seq_to_col_row(state, self.num_cols)
            if self.obs_states is not None:
                end_states = np.vstack((self.obs_states, self.goal_states))
            else:
                end_states = self.goal_states

            if any(np.sum(np.abs(end_states-row_col), 1) == 0):
                self.P[state, state, action] = 1

            # else consider stochastic effects of action
            else:
                for dir in range(-1, 2, 1):

                    direction = self._get_direction(action, dir)
                    next_state = self._get_state(state, direction)
                    if dir == 0:
                        prob = self.p_good_trans
                    elif dir == -1:
                        prob = (1 - self.p_good_trans)*(self.bias)
                    elif dir == 1:
                        prob = (1 - self.p_good_trans)*(1-self.bias)

                    self.P[state, next_state, action] += prob

            # make restart states transition back to the start state with
            # probability 1
            if self.restart_states is not None:

```

```

        if any(np.sum(np.abs(self.restart_states-row_col), 1)==0):
            next_state = GridWorld.row_col_to_seq(self.start_state,
↪self.num_cols)

            self.P[state, :, :] = 0
            self.P[state, next_state, :] = 1

        return self

def _get_direction(self, action, direction):

    left = [2, 3, 1, 0]
    right = [3, 2, 0, 1]
    if direction == 0:
        new_direction = action
    elif direction == -1:
        new_direction = left[action]
    elif direction == 1:
        new_direction = right[action]
    else:
        raise Exception("getDir received an unspecified case")
    return new_direction

def _get_state(self, state, direction):

    row_change = [-1, 1, 0, 0]
    col_change = [0, 0, -1, 1]
    row_col = GridWorld.seq_to_col_row(state, self.num_cols)
    row_col[0, 0] += row_change[direction]
    row_col[0, 1] += col_change[direction]

    # check for invalid states
    if self.obs_states is not None:
        if (np.any(row_col < 0) or
            np.any(row_col[:, 0] > self.num_rows-1) or
            np.any(row_col[:, 1] > self.num_cols-1) or
            np.any(np.sum(abs(self.obs_states - row_col), 1)==0)):
            next_state = state
        else:
            next_state = GridWorld.row_col_to_seq(row_col, self.num_cols)[0]
    else:
        if (np.any(row_col < 0) or
            np.any(row_col[:, 0] > self.num_rows-1) or
            np.any(row_col[:, 1] > self.num_cols-1)):
            next_state = state
        else:
            next_state = GridWorld.row_col_to_seq(row_col, self.num_cols)[0]

    return next_state

```

```

def plot(self):
    """
    Plots the gridworld with the start, goal, and bad states.
    Mark X for Obstructions, G for Goal, B for Bad, and S for Start
    """
    grid = np.zeros((self.num_rows, self.num_cols))
    if self.obs_states is not None:
        for i in range(self.obs_states.shape[0]):
            grid[self.obs_states[i, 0], self.obs_states[i, 1]] = 1
    for i in range(self.goal_states.shape[0]):
        grid[self.goal_states[i, 0], self.goal_states[i, 1]] = -1
    for i in range(self.bad_states.shape[0]):
        grid[self.bad_states[i, 0], self.bad_states[i, 1]] = 2
    for i in range(self.restart_states.shape[0]):
        grid[self.restart_states[i, 0], self.restart_states[i, 1]] = 3
    grid[self.start_state[0, 0], self.start_state[0, 1]] = -2
    sns.heatmap(grid, annot=False, cmap="coolwarm", cbar=False,
    ↪linewidths=0.5, linecolor='black')
    plt.gca().set_aspect('equal', adjustable='box')

    for i in range(self.num_rows):
        for j in range(self.num_cols):
            if grid[i, j] == -1:
                plt.text(j+0.5, i+0.5, 'G', ha='center', va='center',
    ↪fontsize=10)
            if grid[i, j] == 2:
                plt.text(j+0.5, i+0.5, 'B', ha='center', va='center',
    ↪fontsize=10)
            if grid[i, j] == 3:
                plt.text(j+0.5, i+0.5, 'R', ha='center', va='center',
    ↪fontsize=10)
            if grid[i, j] == -2:
                plt.text(j+0.5, i+0.5, 'S', ha='center', va='center',
    ↪fontsize=10)
            if grid[i, j] == 1:
                plt.text(j+0.5, i+0.5, 'X', ha='center', va='center',
    ↪fontsize=10)

def reset(self):
    return int(self.start_state_seq)

def step(self, state, action):
    p, r = 0, np.random.random()
    for next_state in range(self.num_states):

```

```

        p += self.P[state, next_state, action]

        if r <= p:
            break

    if (self.wind and np.random.random() < 0.4):
        arr = self.P[next_state, :, 3]
        next_next = np.where(arr == np.amax(arr))
        next_next = next_next[0][0]
        return next_next, self.R[next_next]
    else:
        return next_state, self.R[next_state]

    @staticmethod
    def row_col_to_seq(row_col, num_cols):
        #Converts state number to row_column format
        return row_col[:, 0] * num_cols + row_col[:, 1]

    @staticmethod
    def seq_to_col_row(seq, num_cols):
        #Converts row_column format to state number
        r = floor(seq / num_cols)
        c = seq - r * num_cols
        return np.array([[r, c]])

```

```

[5]: def get_env(state, wind=False, p_good_transition=1.0):
    gw = GridWorld(
        num_rows=num_rows,
        num_cols=num_columns,
        start_state=start_state[state],
        goal_states=goal_states,
        wind=wind
    )
    gw.add_obstructions(
        obstructed_states=obstructions,
        bad_states=bad_states,
        restart_states=restart_states
    )
    gw.add_transition_probability(
        p_good_transition=p_good_transition,
        bias=bias
    )
    gw.add_rewards(
        step_reward=step_reward,
        goal_reward=goal_reward,
        bad_state_reward=bad_state_reward,

```

```

        restart_state_reward=restart_state_reward
    )
    env = gw.create_gridworld()
    return env

```

## 3 Policy

### 3.1 Action Policy Definitions

```

[6]: class BasePolicy:
    @property
    def name(self):
        raise NotImplementedError

    def select_action(self, state, action_values):
        raise NotImplementedError

```

#### 3.1.1 Greedy Policy

```

[7]: class GreedyPolicy(BasePolicy):
    @property
    def name(self):
        return 'greedy'

    def __init__(self, actions):
        self.actions = actions

    def select_action(self, state, action_values):
        return self.actions[np.argmax(action_values[state, :])]

```

#### 3.1.2 $\varepsilon$ -Greedy Policy

The  $\varepsilon$ -greedy policy defined as

$$\text{next\_action} = \begin{cases} \arg \max_{a \in A(s)} Q(s, a) & \text{with probability } 1 - \varepsilon \\ \text{random choice} & \text{with probability } \varepsilon \end{cases}$$

Hyperparameters: -  $\varepsilon$ : The probability of choosing a random action

```

[8]: class EpGreedyPolicy(BasePolicy):
    @property
    def name(self):
        return f'ep-greedy ep:{self.epsilon}'

    def __init__(self, epsilon, actions):
        self.epsilon = epsilon
        self.actions = actions

```



```

def select_action(self, state, action_values):

    if np.random.binomial(1, 1-self.epsilon):
        return self.actions[np.argmax(action_values[state, :])]
    else:
        return np.random.choice(self.actions)

```

### 3.1.3 Softmax Policy

The softmax policy is defined as

$$\text{next\_action} = \begin{cases} a_1 & \text{with probability } \mathcal{P}(1) \\ a_2 & \text{with probability } \mathcal{P}(2) \\ \vdots & \vdots \\ a_n & \text{with probability } \mathcal{P}(n) \end{cases}$$

where

$$\mathcal{P}(a) = \frac{e^{Q(s,a)/\tau}}{\sum_{i=1}^n e^{Q(s,i)/\tau}}$$

Hyperparameters: -  $\tau$ : The temperature parameter

```

[9]: from scipy.special import softmax

class SoftmaxPolicy(BasePolicy):
    @property
    def name(self):
        return f'softmax tau:{self.tau}'

    def __init__(self, tau, actions):
        self.tau = tau
        self.actions = actions

    def select_action(self, state, action_values):
        return np.random.choice(self.actions, p = softmax(action_values[state, :
↪]/self.tau))

```

## 3.2 Update Policy Definitions

```

[10]: class BaseUpdate:
        @property
        def name(self):
            raise NotImplementedError

        def update(self, state):
            raise NotImplementedError

```

### 3.2.1 SARSA

The update rule for SARSA:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

Hyperparameters: -  $\alpha$ : The learning rate -  $\gamma$ : The discount factor

```
[11]: class SARSAUpdate(BaseUpdate):
    @property
    def name(self):
        return 'sarsa'

    def __init__(self, alpha, gamma):
        self.alpha = alpha
        self.gamma = gamma

    def update(self, Q, state, action, next_state, next_action, reward):
        return Q[state, action] + self.alpha * (reward + self.gamma *
        ↪Q[next_state, next_action] - Q[state, action])
```

### 3.2.2 Q-Learning

The update rule for Q-Learning:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$

Hyperparameters: -  $\alpha$ : The learning rate -  $\gamma$ : The discount factor

```
[12]: class QLearningUpdate(BaseUpdate):
    @property
    def name(self):
        return 'q-learning'

    def __init__(self, alpha, gamma):
        self.alpha = alpha
        self.gamma = gamma

    def update(self, Q, state, action, next_state, next_action, reward):
        return Q[state, action] + self.alpha * (reward + self.gamma * np.
        ↪max(Q[next_state, :]) - Q[state, action])
```

## 3.3 Policy Iterators

```
[13]: class PolicyTrainer:
    def __init__(self, env, exploration_policy, update_policy, episodes, runs):
        self.env = env
        self.exploration_policy = exploration_policy
        self.update_policy = update_policy
```

```

self.episodes = episodes
self.runs = runs

self.steps = np.zeros((runs, episodes))
self.rewards = np.zeros((runs, episodes))
self.Q = np.zeros((runs, env.num_states, env.num_actions))
self.hmap_visits = np.zeros((runs, env.num_states))
self.hmap_Q = np.zeros((runs, env.num_states))

def train(self):
    np.random.seed(32+152)

    for run in tqdm.trange(self.runs, desc='Training Runs'):
        for episode in range(self.episodes):

            current_state = GridWorld.row_col_to_seq(self.env.start_state,
↪self.env.num_cols)[0]
            current_action = self.exploration_policy.
↪select_action(current_state, self.Q[run])

            self.steps[run, episode] = 0
            self.rewards[run, episode] = 0
            self.hmap_visits[run, current_state] += 1

            while current_state not in GridWorld.row_col_to_seq(self.env.
↪goal_states, self.env.num_cols) and self.steps[run, episode] < 100:

                next_state, reward = self.env.step(current_state,
↪current_action)
                next_action = self.exploration_policy.
↪select_action(next_state, self.Q[run])

                self.Q[run, current_state, current_action] = self.
↪update_policy.update(self.Q[run], current_state, current_action, next_state,
↪next_action, reward)

                if current_state != next_state:
                    self.hmap_visits[run, next_state] += 1

                current_state = next_state
                current_action = next_action

            self.steps[run, episode] += 1
            self.rewards[run, episode] += reward

```

```

        if current_state not in list(GridWorld.row_col_to_seq(self.env.
↪goal_states, self.env.num_cols)):
            # self.steps[run, episode] = np.inf
            pass

        for state in range(self.env.num_states):
            self.hmap_Q[run, state] = np.max(self.Q[run, state, :])

    def evaluate(self):
        Q = np.mean(self.Q, axis=0)
        policy = GreedyPolicy(np.arange(self.env.num_actions))
        hmap_visits = np.zeros(self.env.num_states)
        hmap_visits[GridWorld.row_col_to_seq(self.env.start_state, self.env.
↪num_cols)] = 1

        current_state = GridWorld.row_col_to_seq(self.env.start_state, self.env.
↪num_cols)[0]
        current_action = policy.select_action(current_state, Q)

        steps = 0
        rewards = 0

        path = [current_state]
        while current_state not in GridWorld.row_col_to_seq(self.env.
↪goal_states, self.env.num_cols) and steps < 100:

            next_state, reward = self.env.step(current_state, ↵
↪current_action)
            next_action = policy.select_action(next_state, Q)

            hmap_visits[next_state] = 1

            current_state = next_state
            current_action = next_action

            steps += 1
            rewards += reward
            path.append(current_state)

        return rewards, steps, path

    def plot_policy(self):

        start_state = '_s1' if (self.env.start_state == np.array([0, 4])).all() ↵
↪else '_s2'

```

```

        wind_state = '_windy' if self.env.wind else '_clear'
        p_observation = str(self.env.p_good_trans)
        name = self.update_policy.name + start_state + wind_state + "p:" +
        ↪p_observation + self.exploration_policy.name + '_policy.pdf'

        _, _, path = self.evaluate()

        plt.title("Learnt Policy")
        # hmap = sns.heatmap(hmap_visits.reshape(self.env.num_rows, self.env.
        ↪num_cols), annot=False)
        self.env.plot()
        plt.plot([x % self.env.num_cols + 0.5 for x in path], [x // self.env.
        ↪num_cols + 0.5 for x in path], 'r--')
        plt.savefig('plots/' + name, pad_inches=0.1, bbox_inches='tight')
        plt.close()

    def plot_reward(self):

        start_state = '_s1' if (self.env.start_state == np.array([0, 4])).all()
        ↪else '_s2'
        wind_state = '_windy' if self.env.wind else '_clear'
        p_observation = str(self.env.p_good_trans)
        name = self.update_policy.name + start_state + wind_state + "p:" +
        ↪p_observation + self.exploration_policy.name + '_reward.pdf'

        plt.title(f"Reward per Episode: Avg:{round(np.mean(self.rewards), 3)},
        ↪Max:{round(np.max(self.rewards), 3)}")

        plt.plot(self.rewards.mean(axis=0), 'r')
        plt.fill_between(range(self.episodes), self.rewards.mean(axis=0) - self.
        ↪rewards.std(axis=0), self.rewards.mean(axis=0) + self.rewards.std(axis=0),
        ↪alpha=0.2, color='r')

        plt.xlabel('Episodes')
        plt.ylabel('Rewards')
        plt.savefig('plots/' + name, pad_inches=0.1, bbox_inches='tight')
        plt.close()

    def plot_steps(self):

        start_state = '_s1' if (self.env.start_state == np.array([0, 4])).all()
        ↪else '_s2'
        wind_state = '_windy' if self.env.wind else '_clear'
        p_observation = str(self.env.p_good_trans)

```

```

        name = self.update_policy.name + start_state + wind_state + "p:" + \
        ↪p_observation + self.exploration_policy.name + '_steps.pdf'

        plt.title(f"Steps per Episode: Avg:{round(np.mean(self.steps), 3)}, Max:
        ↪{round(np.max(self.steps), 3)}, Min:{round(np.min(self.steps), 3)}")

        plt.plot(self.steps.mean(axis=0), 'b')
        plt.fill_between(range(self.episodes), self.steps.mean(axis=0) - self.
        ↪steps.std(axis=0), self.steps.mean(axis=0) + self.steps.std(axis=0), alpha=0.
        ↪2, color='b')

        plt.xlabel('Episodes')
        plt.ylabel('Steps')
        plt.savefig('plots/' + name, pad_inches=0.1, bbox_inches='tight')
        plt.close()

    def plot_visits(self):

        start_state = '_s1' if (self.env.start_state == np.array([0, 4])).all() \
        ↪else '_s2'
        wind_state = '_windy' if self.env.wind else '_clear'
        p_observation = str(self.env.p_good_trans)
        name = self.update_policy.name + start_state + wind_state + "p:" + \
        ↪p_observation + self.exploration_policy.name + '_heatmap.pdf'

        plt.title("State Visits")
        hmap_visits = np.mean(self.hmap_visits, axis=0)
        hmap = sns.heatmap(hmap_visits.reshape(self.env.num_rows, self.env.
        ↪num_cols), annot=False)
        plt.gca().set_aspect('equal', adjustable='box')
        plt.savefig('plots/' + name, pad_inches=0.1, bbox_inches='tight')
        plt.close()

    def plot_Q(self):

        start_state = '_s1' if (self.env.start_state == np.array([0, 4])).all() \
        ↪else '_s2'
        wind_state = '_windy' if self.env.wind else '_clear'
        p_observation = str(self.env.p_good_trans)
        name = self.update_policy.name + start_state + wind_state + "p:" + \
        ↪p_observation + self.exploration_policy.name + '_Q.pdf'

        plt.title(f"Q Value: Avg:{round(np.mean(self.hmap_Q), 3)}, Max:
        ↪{round(np.max(self.hmap_Q), 3)}, Min:{round(np.min(self.hmap_Q), 3)}")
        Q = np.mean(self.hmap_Q, axis=0)

```

```

hmap = sns.heatmap(Q.reshape(self.env.num_rows, self.env.num_cols),
↳annot=False, linewidths=0.5, linecolor='black', cmap="plasma")

for i in range(self.env.num_rows):
    for j in range(self.env.num_cols):
        if [i, j] in self.env.obs_states.tolist():
            plt.text(j+0.5, i+0.5, 'X', ha='center', va='center',
↳fontsize=10)
            continue
        if [i, j] in self.env.goal_states.tolist():
            plt.text(j+0.5, i+0.5, 'G', ha='center', va='center',
↳fontsize=10)
            continue
        if [i, j] in self.env.bad_states.tolist():
            plt.text(j+0.5, i+0.5, 'B', ha='center', va='center',
↳fontsize=10)
            continue
        if [i, j] in self.env.restart_states.tolist():
            plt.text(j+0.5, i+0.5, 'R', ha='center', va='center',
↳fontsize=10)
            continue
        if [i, j] in self.env.start_state.tolist():
            plt.text(j+0.5, i+0.5, 'S', ha='center', va='center',
↳fontsize=10, color='white')

state = GridWorld.row_col_to_seq(np.array([[i, j]]), self.env.
↳num_cols)[0]
action = np.argmax(self.Q[0, state, :])

if action == 0:
    plt.arrow(j+0.5, i+0.5+0.2, 0, -0.4, head_width=0.1,
↳head_length=0.1, fc='k', ec='k')
elif action == 1:
    plt.arrow(j+0.5, i+0.5-0.2, 0, 0.4, head_width=0.1,
↳head_length=0.1, fc='k', ec='k')
elif action == 2:
    plt.arrow(j+0.5+0.2, i+0.5, -0.4, 0, head_width=0.1,
↳head_length=0.1, fc='k', ec='k')
elif action == 3:
    plt.arrow(j+0.5-0.2, i+0.5, 0.4, 0, head_width=0.1,
↳head_length=0.1, fc='k', ec='k')

plt.gca().set_aspect('equal', adjustable='box')
plt.savefig('plots/' + name, pad_inches=0.1, bbox_inches='tight')
plt.close()

```

## 4 Hyperparameter Tuning

To conduct hyperparameter tuning, we opt for maximizing asymptotic optimality, which entails leveraging Q-values acquired by the agent and employing a greedy action selection method.

Following this approach, we establish a grid search function to determine the optimal hyperparameter set based on asymptotic optimality.

```
[14]: def reward_grid_search(env, alphas, gammas, epsilons, taus, model = 'sarsa',  
    ↪ policy = 'epsilon'):  
    optimal_reward = - np.inf  
    best_reward = - np.inf  
    optimal_hyperparams = {}  
  
    if policy == "softmax":  
        # Softmax  
        for gamma in gammas:  
            for alpha in alphas:  
                for tau in taus:  
                    print(f"The current set of Hyperparams: alpha = {alpha},  
    ↪ gamma = {gamma}, tau = {tau}")  
                    if model == "sarsa":  
                        update_policy = SARSAUpdate(alpha=alpha, gamma=gamma)  
                    elif model == "qlearning":  
                        update_policy = QLearningUpdate(alpha=alpha,  
    ↪ gamma=gamma)  
  
                    exploration_policy = SoftmaxPolicy(tau=tau, actions=np.  
    ↪ arange(env.num_actions))  
                    trainer = PolicyTrainer(env, exploration_policy,  
    ↪ update_policy, episodes=10000, runs=5)  
                    trainer.train()  
                    greedy_reward, _, _ = trainer.evaluate()  
                    reward = trainer.rewards  
                    mean_reward = np.mean(np.mean(reward, axis = 1), axis = 0)  
                    if optimal_reward < mean_reward and best_reward <   
    ↪ greedy_reward:  
                        optimal_trainer = trainer  
                        best_reward = greedy_reward  
                        optimal_reward = mean_reward  
                        optimal_hyperparams = {  
                            "alpha": alpha,  
                            "gamma": gamma,  
                            "tau": tau  
                        }  
  
    elif policy == "epsilon":  
        # Epsilon
```



```

        for gamma in gammas:
            for alpha in alphas:
                for epsilon in epsilons:
                    print(f"The current set of Hyperparams: alpha = {alpha},  

↪gamma = {gamma}, epsilon = {epsilon}")
                    if model == "sarsa":
                        update_policy = SARSAUpdate(alpha=alpha, gamma=gamma)
                    elif model == "qlearning":
                        update_policy = QLearningUpdate(alpha=alpha,  

↪gamma=gamma)

                    exploration_policy = EpGreedyPolicy(epsilon=epsilon,  

↪actions=np.arange(env.num_actions))
                    trainer = PolicyTrainer(env, exploration_policy,  

↪update_policy, episodes=10000, runs=5)
                    trainer.train()
                    greedy_reward, _, _ = trainer.evaluate()
                    reward = trainer.rewards
                    mean_reward = np.mean(np.mean(reward, axis = 1), axis = 0)
                    if optimal_reward < mean_reward and best_reward <   

↪greedy_reward:

                        optimal_trainer = trainer
                        best_reward = greedy_reward
                        optimal_reward = mean_reward
                        optimal_hyperparams = {
                            "alpha": alpha,
                            "gamma": gamma,
                            "epsilon": epsilon
                        }

    return optimal_hyperparams, optimal_trainer

```

$\alpha$ : 0.001, 0.01, 0.1, 0.2

$\gamma$ : 0.7, 0.8, 0.9, 1

$\epsilon$ : 0.001, 0.01, 0.1, 0.5

$\tau$ : 0.01, 0.1, 1, 2

```

[15]: # alphas = [0.01, 0.05, 0.1]
      # gammas = [0.8, 0.9, 1]
      # epsilons = [0.001, 0.01, 0.1]
      # tau = [0.01, 0.1, 1]

      alphas = [0.1]
      gammas = [0.8, 0.9]
      epsilons = [0.01]
      tau = [0.01, 0.1]

```

#### 4.0.1 Experiment 1

State = s1 (0, 4) Wind = False p = 1 SARSA Algo

```
[16]: # env = get_env('s1', wind=False, p_good_transition=1.0)
# update_policy = SARSAUpdate(alpha=0.1, gamma=0.8)
# exploration_policy = EpGreedyPolicy(epsilon=0.2, actions=np.arange(env.
    ↪ num_actions))
# trainer = PolicyTrainer(env, exploration_policy, update_policy,
    ↪ episodes=10000, runs=5)
# trainer.train()
# trainer.plot_reward()
# trainer.plot_steps()
# trainer.plot_visits()
# trainer.plot_Q()
# trainer.plot_policy()

[17]: env = get_env('s1', wind=False, p_good_transition=1.0)

# Softmax
optimal_hyperparams, optimal_trainer = reward_grid_search(env, alphas, gammas,
    ↪ epsilons, tau, model = "sarsa", policy = "softmax")
print(f"Optimal Hyperparameters for Softmax - alpha:
    ↪ {optimal_hyperparams['alpha']}, gamma: {optimal_hyperparams['gamma']}, tau:
    ↪ {optimal_hyperparams['tau']}")
# Plotting
optimal_trainer.plot_reward()
optimal_trainer.plot_steps()
optimal_trainer.plot_visits()
optimal_trainer.plot_Q()
optimal_trainer.plot_policy()

# Epsilon Greedy
optimal_hyperparams, optimal_trainer = reward_grid_search(env, alphas, gammas,
    ↪ epsilons, tau, model = "sarsa", policy = "epsilon")
print(f"Optimal Hyperparameters for Epsilon Greedy - alpha:
    ↪ {optimal_hyperparams['alpha']}, gamma: {optimal_hyperparams['gamma']},
    ↪ epsilon: {optimal_hyperparams['epsilon']}")
# Plotting
optimal_trainer.plot_reward()
optimal_trainer.plot_steps()
optimal_trainer.plot_visits()
optimal_trainer.plot_Q()
optimal_trainer.plot_policy()
```

The current set of Hyperparams: alpha = 0.1, gamma = 0.8, tau = 0.01

Training Runs: 100%|  
| 5/5 [00:25<00:00, 5.12s/it]

The current set of Hyperparams:  $\alpha = 0.1$ ,  $\gamma = 0.8$ ,  $\tau = 0.1$

Training Runs: 100%|  
| 5/5 [00:19<00:00, 3.98s/it]

The current set of Hyperparams:  $\alpha = 0.1$ ,  $\gamma = 0.9$ ,  $\tau = 0.01$

Training Runs: 100%|  
| 5/5 [00:17<00:00, 3.54s/it]

The current set of Hyperparams:  $\alpha = 0.1$ ,  $\gamma = 0.9$ ,  $\tau = 0.1$

Training Runs: 100%|  
| 5/5 [00:21<00:00, 4.30s/it]

Optimal Hyperparameters for Softmax -  $\alpha: 0.1$ ,  $\gamma: 0.8$ ,  $\tau: 0.01$

The current set of Hyperparams:  $\alpha = 0.1$ ,  $\gamma = 0.8$ ,  $\epsilon = 0.01$

Training Runs: 100%|  
| 5/5 [00:08<00:00, 1.65s/it]

The current set of Hyperparams:  $\alpha = 0.1$ ,  $\gamma = 0.9$ ,  $\epsilon = 0.01$

Training Runs: 100%|  
| 5/5 [00:07<00:00, 1.54s/it]

Optimal Hyperparameters for Epsilon Greedy -  $\alpha: 0.1$ ,  $\gamma: 0.8$ ,  $\epsilon: 0.01$

[ ]: